

# OS Project1 Final Report

letsgo\_seuljongbin

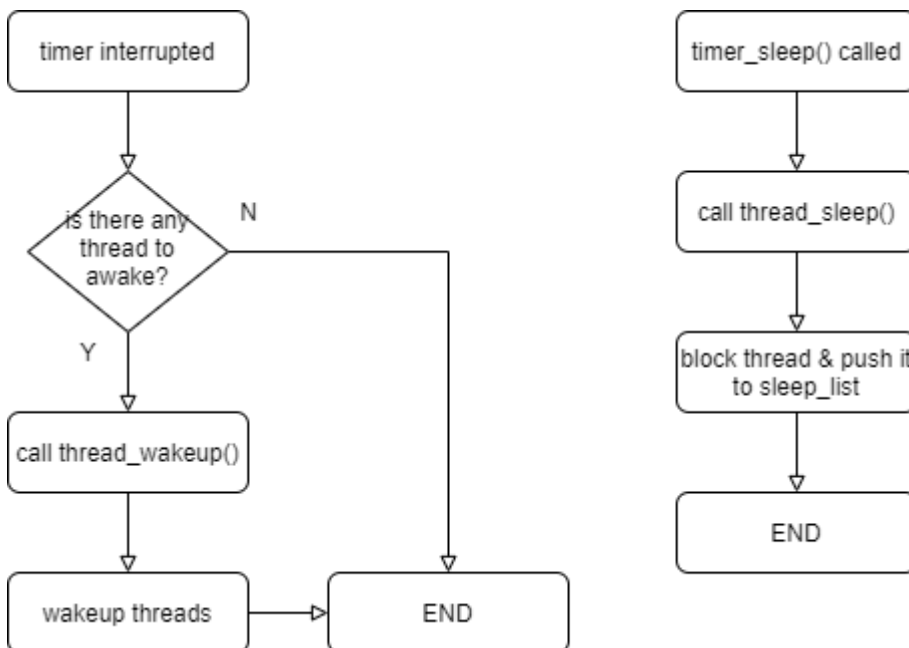
20180038 박형규, 20180480 성장환

## 1. Introduction

OS Project1을 통해 Pintos Alarm Clock, Priority Scheduling, Advanced Scheduler를 모두 구현해 보았다. 이번 Final Report에서는 각 항목별로 Flow Chart, 추가하거나 수정한 Function & Data Structure, Implementation에 대해 자세히 작성하였다. 특히 Implementation 부분에서는 Design Report와 비교하여 작성하였는데 대부분 Design Report와 동일하지만 Design Report에서 언급되지 않았거나 바뀐 부분이 있다면 해당 부분을 언급하였다.

## 2. Alarm Clock

### 2.1. Flow Chart



### 2.2. Added or modified Function & Data Structure

- Modified Data Structure  
struct thread, static int64\_t first\_time\_to\_wakeup, static struct list sleep\_list
- Added Function  
get\_first\_time\_to\_wakeup(), thread\_sleep(), thread\_wakeup()
- Modified Function  
thread\_init(), timer\_sleep(), timer\_interrupt()

## 2.3. Implementation

우선 struct thread에서 thread를 깨워야 할 tick을 저장할 변수 time\_to\_wakeup을 선언했다.

```
93     int64_t time_to_wakeup;
```

다음으로 sleep상태의 thread들을 관리할 sleep\_list를 선언하고 thread\_init()에서 초기화 시켰다.

```
32     static struct list sleep_list; 100     list_init(&sleep_list);
```

다음으로 sleep\_list에 있는 thread들 중 가장 먼저 깨워야 하는 thread가 일어나야 할 시각을 저장하는 전역변수 first\_time\_to\_wakeup을 선언했다.

```
34     static int64_t first_time_to_wakeup;
```

다음으로 thread를 sleep시키는 함수 thread\_sleep()을 구현했다. thread\_sleep()은 인자로 현재 tick(start)과 몇 tick동안 sleep시킬지(ticks)를 입력 받아서 current thread를 sleep시킨다. 이때, interrupt를 차단하고, time\_to\_wakeup과 first\_time\_to\_wakeup에 언제 이 thread가 wakeup될지 저장한 다음 sleep\_list에 push back해준다. 그런 다음 thread\_block()을 해서 BLOCKED로 설정한다.

```
115 void thread_sleep(int64_t start, int64_t ticks)
116 {
117     struct thread *cur;
118     enum intr_level old_level;
119     old_level = intr_disable();
120
121     ASSERT(thread_current() != idle_thread);
122
123     cur = thread_current();
124     cur->time_to_wakeup = start + ticks;
125
126     if(cur->time_to_wakeup < first_time_to_wakeup) first_time_to_wakeup = cur->time_to_wakeup;
127
128     list_push_back(&sleep_list, &cur->elem);
129
130     thread_block();
131
132     intr_set_level(old_level);
133 }
```

다음으로 sleep상태의 thread를 깨우는 함수 thread\_wakeup을 구현했다. Sleep\_list를 처음부터 끝까지 순회하면서 깨워야되는 thread들을 sleep\_list에서 제거하고 unblock해준다. 아직 깨울 시간이 되지 않은 thread들은 넘어가고 first\_time\_to\_wakeup을 다음에 깨워야하는 tick으로 설정하도록 구현했다.

```
135 void thread_wakeup(int64_t ticks)
136 {
137     first_time_to_wakeup = INT64_MAX;
138     struct list_elem *e;
139     struct thread *t;
140
141     for(e = list_begin(&sleep_list); e != list_end(&sleep_list);)
142     {
143         t = list_entry(e, struct thread, elem);
144
145         if(ticks < t->time_to_wakeup)
146         {
147             e = list_next(e);
148             if(t->time_to_wakeup < first_time_to_wakeup) first_time_to_wakeup = t->time_to_wakeup;
149         }
150         else
151         {
152             e = list_remove(&t->elem);
153             thread_unblock(t);
154         }
155     }
156 }
```

다음으로 thread\_wakeup이 실행되게 하기 위해 timer\_interrupt를 수정했다.

```
184     if((get_first_time_to_wakeup() <= ticks))
185     |   thread_wakeup(ticks);
```

이때, 위와 같이 first\_time\_to\_wakeup 값을 현재 tick과 비교하기 위해서 first\_time\_to\_wakeup을 리턴해주는 get\_first\_time\_to\_wakeup()을 추가로 구현해 주었다.

```
110 int64_t get_first_time_to_wakeup(void)
111 {
112     return first_time_to_wakeup;
113 }
```

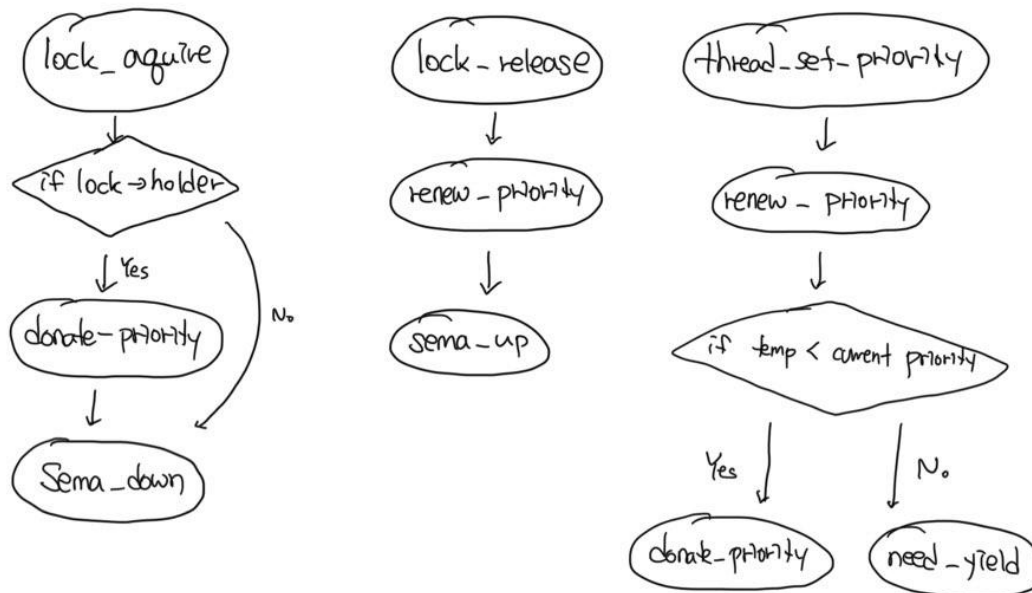
마지막으로 기본 pintos의 timer\_sleep()에 구현되어 있던 busy waiting방식 코드를 제거하고, thread\_sleep()을 호출하도록 수정했다.

```
90 timer_sleep (int64_t ticks)
91 {
92     int64_t start = timer_ticks ();
93
94     ASSERT (intr_get_level () == INTR_ON);
95
96     thread_sleep(start, ticks);
97 }
```

### 3. Priority Scheduling

#### 3.1. Flow Chart

< Priority Donation Flow Chart >



#### 3.2. Added or modified Function & Data Structure

- Modified Data Structure

struct thread

- Added Function

priority\_compare(), need\_yield(), sema\_priority\_compare(), donate\_priority(), renew\_priority()

- Modified Function

thread\_yield(), thread\_unblock(), thread\_create(), thread\_set\_priority(), sema\_down(), sema\_up(),  
cond\_wait(), cond\_signal(), lock\_acquire(), lock\_release(), thread\_set\_priority(), thread\_get\_priority()

#### 3.3. Implementation

- Priority Scheduling & Synchronization

Priority Scheduling은 Design Report에서 구현한 방법과 같이 구현하였다. 기존의 방법에서 thread\_yield() 함수와 thread\_unblock() 함수에서 ready\_list에 thread를 insert할 때 사용하던 방법인 list\_push\_back() 대신에 list\_insert\_ordered() 함수를 사용하여 ready\_list 내부에서 thread의 우선순위를 기준으로 정렬되도록 하였다.

```

void
thread_yield (void)
{
    struct thread *cur = thread_current ();
    enum intr_level old_level;

    ASSERT (!intr_context ());

    old_level = intr_disable ();
    if (cur != idle_thread)
        list_insert_ordered (&ready_list, &cur->elem, priority_compare, 0);
    cur->status = THREAD_READY;
    schedule ();
    intr_set_level (old_level);
}

```

```

void
thread_unblock (struct thread *t)
{
    enum intr_level old_level;

    ASSERT (is_thread (t));

    old_level = intr_disable ();
    ASSERT (t->status == THREAD_BLOCKED);
    list_insert_ordered (&ready_list, &t->elem, priority_compare, 0);
    t->status = THREAD_READY;
    intr_set_level (old_level);
}

```

이 때, list\_insert\_ordered() 함수의 인자로 list에 ordered하게 element를 집어넣어야 하므로 thread간의 priority를 비교하는 함수가 필요하다. 따라서 priority\_compare() 함수를 구현하였다.

```

bool priority_compare(const struct list_elem *a, const struct list_elem *b, void *aux UNUSED)
{
    const struct thread *a_ = list_entry (a, struct thread, elem);
    const struct thread *b_ = list_entry (b, struct thread, elem);

    if(a->priority > b->priority)
        return 1;
    else
        return 0;
}

```

다음으로 ready\_list에서 가장 priority가 큰 thread(ready\_list의 첫 번째 thread)보다 현재의 thread보다 priority가 높으면 양보를 하도록 해야 하는 부분을 thread\_create() 부분의 마지막에 넣어야 한다고 Design Report에 언급하였다. 우리는 이 함수를 need\_yield() 함수로 따로 만들어서 구현하였다. thread\_create() 함수 아래에 이 부분을 써넣지 않고 따로 만든 이유는 이 부분이 thread\_set\_priority() 함수 내부에도 필요하기 때문이다. 이 부분은 Design Report의 Priority Donate부분에서 마지막 켄에 '이 작업에서 current thread가 더이상 highest priority를 가지고 있지 않다면(다면-오타) thread\_yield()함수를 호출하는 작업이 필요하다.' 라고 언급되어 있다. 이 부분이 priority scheduling만을 구현하였을 때는 510, 511번째 줄만 구현되어 있었는데 이후에 Priority Donation을 구현하게 되면서 위 부분을 추가하게 된다. 이는 2.3에서 이어 설명하도록 하겠다.

```

325 tid_t
326 thread_create (const char *name, int priority,
327               thread_func *function, void *aux)
328 {
329     struct thread *t;
330     struct kernel_thread_frame *kf;
331     struct switch_entry_frame *ef;
332     struct switch_threads_frame *sf;
333     tid_t tid;
334
335     ASSERT (function != NULL);
336
337     /* Allocate thread. */
338     t = palloc_get_page (PAL_ZERO);
339     if (t == NULL)
340         return TID_ERROR;
341
342     /* Initialize thread. */
343     init_thread (t, name, priority);
344     tid = t->tid = allocate_tid ();
345
346     /* Stack frame for kernel_thread(). */
347     kf = alloc_frame (t, sizeof *kf);
348     kf->eip = NULL;
349     kf->function = function;
350     kf->aux = aux;
351
352     /* Stack frame for switch_entry(). */
353     ef = alloc_frame (t, sizeof *ef);
354     ef->eip = (void (*) (void)) kernel_thread;
355
356     /* Stack frame for switch_threads(). */
357     sf = alloc_frame (t, sizeof *sf);
358     sf->eip = switch_entry;
359     sf->ebp = 0;
360
361     /* Add to run queue. */
362     thread_unblock (t);
363
364     need_yield();
365
366     return tid;
367 }

```

```

497 void
498 thread_set_priority (int new_priority)
499 {
500     if (thread_mlfqs) return;
501     enum intr_level old_level = intr_disable();
502
503     int temp = thread_current()->priority;
504     thread_current()->original_priority = new_priority;
505     renew_priority();
506
507     if(temp < thread_current()->priority)
508         donate_priority();
509
510     if(temp > thread_current()->priority)
511         need_yield();
512
513     intr_set_level(old_level);
514 }

```

다음으로는 Synchronization을 구현한 방법에 대해 설명하도록 하겠다. Design Report에서는 sema\_down() 함수나 sema\_up() 함수 둘 중에 하나만 수정해도 된다고 하였고 우리 구현에서는 sema\_down()을 수정하였다. 먼저 sema\_down() 함수에서는 Design Report에서 언급한 것과 동일하게 waiters list에 element를 insert할 때 list\_push\_back() 함수 대신에 list\_insert\_ordered() 함수를 사용하였다. 아래 사진에서 71번째 줄 donate\_priority()는 이후 priority donation 부분을 설명할 때 언급하도록 하겠다.

```

60 void
61 sema_down (struct semaphore *sema)
62 {
63     enum intr_level old_level;
64
65     ASSERT (sema != NULL);
66     ASSERT (!intr_context ());
67
68     old_level = intr_disable ();
69     while (sema->value == 0)
70     {
71         donate_priority();
72
73         list_insert_ordered(&sema->waiters, &thread_current()->elem, priority_compare, NULL);
74         thread_block ();
75     }
76     sema->value--;
77     intr_set_level (old_level);
78 }

```

sema\_up() 함수를 수정한 부분은 Design Report에 언급되어 있지 않다. 당시 Design Report를 작성할 때에는 미처 생각하지 못한 부분인데, thread가 waiters list에 있는 동안 priority가 변경되었을 경우가 있을 수 있으므로 thread\_unblock()을 수행하기 전에 list\_sort() 함수를 이용하여 waiters list를 priority 기준으로 정렬하도록 하였다. 또한, sema\_up이 하는 역할이 semaphore을 반환하고 동작을 마치므로 thread\_yield()

를 추가해 주었다.

```
110 void
111 sema_up (struct semaphore *sema)
112 {
113     enum intr_level old_level;
114
115     ASSERT (sema != NULL);
116
117     old_level = intr_disable ();
118     if (!list_empty (&sema->waiters))
119     {
120         list_sort(&sema->waiters, priority_compare, NULL);
121         thread_unblock (list_entry(list_pop_front(&sema->waiters), struct thread, elem));
122     }
123     sema->value++;
124
125     thread_yield();
126
127     intr_set_level (old_level);
128 }
```

다음으로는 semaphore와 비슷하게 condition variable과 관련하여 수정한 내용이다. 이 부분도 Design Report를 작성할 때는 구현해야 한다는 것을 미처 생각하지 못하여 Design Report에는 언급되어 있지 않다. 먼저 cond\_wait() 함수에서 condition variable의 waiters list에 element가 priority를 기준으로 삽입되도록 list\_insert\_ordered() 함수를 이용하였다. 이 때, list\_insert\_ordered() 함수의 인자로 sema\_priority\_compare() 함수를 새롭게 만들어 사용하였다. 기존의 priority\_compare() 함수를 사용하지 않고 sema\_priority\_compare() 함수를 따로 만들어 사용한 이유는 condition variable같은 경우 waiters list에 semaphore\_element가 삽입되므로 구조체의 구성이 다르기 때문이다.

```
318 void
319 cond_wait (struct condition *cond, struct lock *lock)
320 {
321     struct semaphore_elem waiter;
322
323     ASSERT (cond != NULL);
324     ASSERT (lock != NULL);
325     ASSERT (!intr_context ());
326     ASSERT (lock_held_by_current_thread (lock));
327
328     sema_init (&waiter.semaphore, 0);
329     list_insert_ordered(&cond->waiters, &waiter.elem, sema_priority_compare, NULL);
330     lock_release (lock);
331     sema_down (&waiter.semaphore);
332     lock_acquire (lock);
333 }
375 bool sema_priority_compare(const struct list_elem *a, const struct list_elem *b, void *aux UNUSED)
376 {
377     struct semaphore_elem *sa = list_entry(a, struct semaphore_elem, elem);
378     struct semaphore_elem *sb = list_entry(b, struct semaphore_elem, elem);
379
380     if(list_empty(&sb->semaphore.waiters))
381         return true;
382
383     if(list_empty(&sa->semaphore.waiters))
384         return false;
385
386     struct thread* ta = list_entry(list_front(&sa->semaphore.waiters), struct thread, elem);
387     struct thread* tb = list_entry(list_front(&sb->semaphore.waiters), struct thread, elem);
388
389     return ta->priority > tb->priority;
390 }
```

마지막으로 cond\_signal() 함수에서 sema\_up()을 하기 전에 waiters list에서 대기하던 중에 priority가 변경되었을 가능성이 있다. 따라서 list\_sort() 함수를 사용하여 재정렬 시켜주었다.

```

342 void
343 cond_signal (struct condition *cond, struct lock *lock UNUSED)
344 {
345     ASSERT (cond != NULL);
346     ASSERT (lock != NULL);
347     ASSERT (!intr_context ());
348     ASSERT (lock_held_by_current_thread (lock));
349
350     if (!list_empty (&cond->waiters)) |
351     {
352         list_sort(&cond->waiters, sema_priority_compare, NULL);
353
354         sema_up (&list_entry (list_pop_front (&cond->waiters),
355         | | | | | | | | struct semaphore_elem, elem)->semaphore);
356     }
357 }

```

#### ■ Priority Donation 구현

먼저 thread 구조체에 priority donation에 필요한 element들을 추가하였다. original\_priority는 donation 이후에 priority를 초기화 시켜 주기 위해 저장한 priority 변경 이전의 priority를 나타내고, want\_lock은 acquire하고 싶은 lock을 나타낸다. donation\_list와 donation\_elem은 multiple donation을 고려하기 위해 도입한 element 들로 donation을 받은 thread 구조체들의 list이다.

```

83 struct thread
84 {
85     /* Owned by thread.c. */
86     tid_t tid; /* Thread identifier. */
87     enum thread_status status; /* Thread state. */
88     char name[16]; /* Name (for debugging purposes). */
89     uint8_t *stack; /* Saved stack pointer. */
90     int priority; /* Priority. */
91     struct list_elem allelem; /* List element for all threads. */
92
93     int64_t time_to_wakeup;
94     int original_priority;
95     int nice;
96     int recent_cpu;
97     struct lock *want_lock;
98     struct list donation_list;
99     struct list_elem donation_elem;
100
101     /* Shared between thread.c and synch.c. */
102     struct list_elem elem; /* List element. */
103
104 #ifdef USERPROG
105     /* Owned by userprog/process.c. */
106     uint32_t *pagedir; /* Page directory. */
107 #endif
108
109     /* Owned by thread.c. */
110     unsigned magic; /* Detects stack overflow. */
111 };
112

```

먼저 init\_thread() 함수에서 thread 구조체에 새롭게 추가한 element들을 초기화 시켜주었다.



```

640 static void
641 init_thread (struct thread *t, const char *name, int priority)
642 {
643     enum intr_level old_level;
644
645     ASSERT (t != NULL);
646     ASSERT (PRI_MIN <= priority && priority <= PRI_MAX);
647     ASSERT (name != NULL);
648
649     memset (t, 0, sizeof *t);
650     t->status = THREAD_BLOCKED;
651     strcpy (t->name, name, sizeof t->name);
652     t->stack = (uint8_t *) t + PGSIZE;
653     t->priority = priority;
654     t->recent_cpu = 0;
655     t->nice = 0;
656     t->magic = THREAD_MAGIC;
657
658     t->original_priority = priority;
659     list_init(&t->donation_list);
660     t->want_lock = NULL;
661
662     old_level = intr_disable ();
663     list_push_back (&all_list, &t->allelem);
664     intr_set_level (old_level);
665 }

```

lock\_acquire() 함수에서는 인자로 받은 lock의 holder가 이미 존재한다면 현재의 thread의 want\_lock 변수에 lock을 저장해 주고, list\_insert\_ordered() 함수를 이용하여 holder의 donation list에 thread\_current의 donation\_elem을 저장하고 donate\_priority() 함수를 호출하여 priority donation을 수행하였다.

```

199 void
200 lock_acquire (struct lock *lock)
201 {
202     ASSERT (lock != NULL);
203     ASSERT (!intr_context ());
204     ASSERT (!lock_held_by_current_thread (lock));
205
206     enum intr_level old_level = intr_disable();
207
208     if(!thread_mlfqs && lock->holder)
209     {
210         thread_current()->want_lock = lock;
211         list_insert_ordered(&lock->holder->donation_list, &thread_current()->donation_elem, priority_compare, NULL);
212         donate_priority();
213     }
214     sema_down (&lock->semaphore);
215     thread_current()->want_lock = NULL;
216     lock->holder = thread_current ();
217
218     intr_set_level(old_level);
219 }

```

다음으로 priority donation이 직접 실행되는 함수인 donate\_priority() 함수이다. 이 함수에서는 처음에 현재 thread가 기다리고 있는 lock에 대해 마치 linked list 형태와 같이 while문을 통해 nested donation을 해결해 준다.

```

180 void donate_priority(void)
181 {
182     // nest donation을 위해서 반복문을 이용해 holder들에게 도네이션 한다.
183     int t = 0;
184     struct thread *thrd = thread_current();
185     struct lock *lock = thrd->want_lock;
186     while (lock != NULL && t < 22){ //임의의 depth
187         t++;
188         if (lock->holder == NULL){
189             return;
190         }
191         if (lock->holder->priority >= thrd->priority){
192             return;
193         }
194         lock->holder->priority = thrd->priority;
195         thrd = lock->holder;
196         lock = thrd->want_lock;
197     }
198 }

```

다음으로 lock release() 에서는 lock을 해지했을 때 이전의 상태로 돌아가기 위해 필요한 작업들을 추가해 주었다. 원래 Design Report에서는 이를 두 개의 함수로 나누어 작업하려고 했었지만, 첫 번째로 만들려고 했던 donation list에서 해지 할 lock을 보유하고 있는 entry를 삭제하는 함수는 lock\_release() 함수에 서밖에 쓰이지 않기 때문에 따로 함수로 만들지 않고 직접 코드를 써 넣었다. 그리고 renew\_priority() 함수에서는 thread의 priority가 변경되었으므로 donation을 고려해서 priority를 다시 결정하는 작업을 진행 하였다. 이는 현재 thread를 donation 받기 전의 original priority로 변경하고, donation list에서 가장 priority가 높은 thread와 현재 thread의 priority를 비교하여 높은 값을 현재 thread의 priority로 설정하는 과정으로 이루어졌다.

```

246 void
247 lock_release (struct lock *lock)
248 {
249     ASSERT (lock != NULL);
250     ASSERT (lock_held_by_current_thread (lock));
251
252     enum intr_level old_level = intr_disable();
253
254     if (!thread_mlfqs){
255         for(struct list_elem *e = list_begin(&thread_current()->donation_list); e != list_end(&thread_current()->donation_list); e = list_next(e)){
256             {
257                 if(list_entry(e, struct thread, donation_elem)->want_lock == lock)
258                     list_remove(e);
259             }
260             renew_priority();
261         }
262     }
263     lock->holder = NULL;
264     sema_up (&lock->semaphore);
265     intr_set_level(old_level);
266 }
267
268 void renew_priority(void)
269 {
270     thread_current()->priority = thread_current()->original_priority;
271
272     if(list_empty(&thread_current()->donation_list))
273         return;
274
275     if(list_entry(list_front(&thread_current()->donation_list), struct thread, donation_elem)->priority > thread_current()->priority)
276         thread_current()->priority = list_entry(list_front(&thread_current()->donation_list), struct thread, donation_elem)->priority;
277 }

```

마지막으로 thread\_set\_priority() 함수와 thread\_get\_priority() 함수의 수정이 이루어졌다. thread\_set\_priority() 함수에서는 현재 thread의 priority를 new\_priority로 설정하려고 할 때 original\_priority를 new priority로 설정하고 renew\_priority() 함수를 이용하여 thread\_current의 priority를 설정해 주었다. 이는 donation을 반영하였기 때문에 나온 결과이다. 이 때 temp < thread\_current()->priority일 경우에는 priority inversion 현상이 발생할 수 있으므로 donate\_priority() 함수를 호출하여 이를 해결한다. thread\_get\_priority() 함수에서는 interrupt disable 과정을 추가해 주어 thread의 priority 값을 return 하도록 하였다.

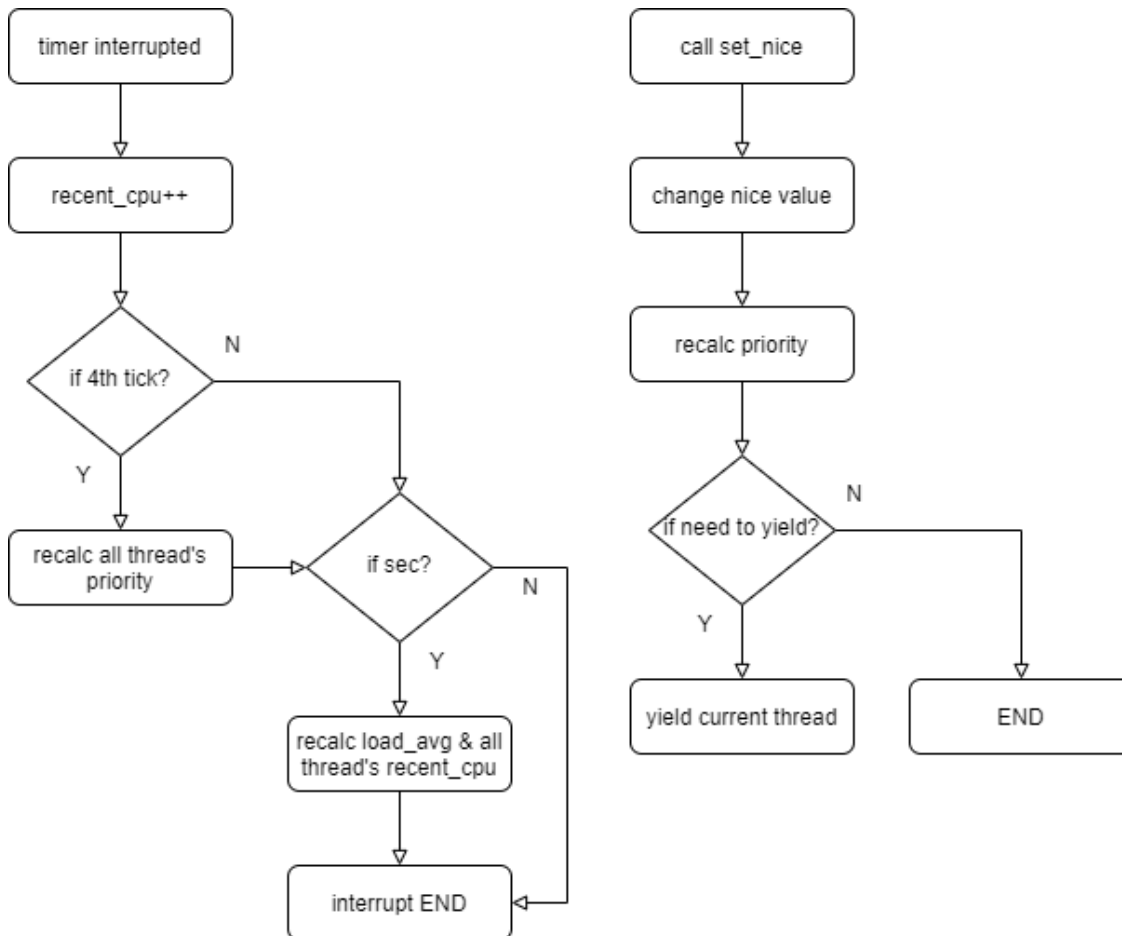
```

487  /* Sets the current thread's priority to NEW_PRIORITY. */
488  void
489  thread_set_priority (int new_priority)
490  {
491      if (thread_mlfqs) return;
492      enum intr_level old_level = intr_disable();
493
494      int temp = thread_current ()->priority;
495      thread_current()->original_priority = new_priority;
496      renew_priority();
497
498      if(temp < thread_current()->priority)
499          donate_priority();
500
501      if(temp > thread_current()->priority)
502          need_yield();
503
504      intr_set_level(old_level);
505  }
506
507  /* Returns the current thread's priority. */
508  int
509  thread_get_priority (void)
510  {
511      enum intr_level old_level = intr_disable();
512      int temp = thread_current()->priority;
513      intr_set_level(old_level);
514      return temp;
515  }

```

## 4. Advanced Scheduler

### 4.1. Flow Chart



### 4.2. Added or modified Function & Data Structure

- Modified Data Structure  
struct thread, int load\_avg
- Added Function  
"fixed\_point.h", mlfqs\_calc\_priority(), mlfqs\_increment\_recent\_cpu(), mlfqs\_recalc\_per\_Sec(), mlfqs\_recalc\_all\_priority()
- Modified Function  
init\_thread(), thread\_start(), thread\_set\_priority(), thread\_set\_nice(), thread\_get\_nice(), thread\_get\_load\_avg(), thread\_get\_recent\_cpu(), timer\_interrupt(), lock\_acquire(), lock\_release()

### 4.3. Implementation

우선 변수 load\_avg를 선언하고 thread\_start()에서 초기화 시켰다.

```
107 /* 구현한 부분 START*/
108 int load_avg; 271 load_avg = 0;
```

다음으로는 struct thread에 변수 nice와 recent\_cpu를 추가하고 init\_thread()에서 초기화 시켰다.

95	int nice;	656	t->recent_cpu = 0;
96	int recent_cpu;	657	t->nice = 0;

다음으로 fixed point 연산을 위한 함수들을 APPENDIX에서 주어진 수식을 따라 "fixed\_point.h"에 다음과 같이 구현했다.

```

1  #define E 17
2  #define F 14
3  #define FRACTION (1<<F)
4
5  int convert_to_fp(int n);
6  int convert_to_int_round_zero(int x);
7  int convert_to_int_round_nearest(int x);
8  int add_fp_fp(int x, int y);
9  int add_fp_int(int x, int n);
10 int sub_fp_fp(int x, int y);
11 int sub_fp_int(int x, int n);
12 int mul_fp_fp(int x, int y);
13 int mul_fp_int(int x, int n);
14 int div_fp_fp(int x, int y);
15 int div_fp_n(int x, int n);
16
17 int convert_to_fp(int n){
18     // Convert n to fixed point
19     return n * (FRACTION);
20 }
21
22 int convert_to_int_round_zero(int x){
23     // Convert x to integer (rounding toward zero)
24     return x / (FRACTION);
25 }
26
27 int convert_to_int_round_nearest(int x){
28     // Convert x to integer (rounding to nearest)
29     if (x >= 0){
30         return (x + (FRACTION)/2)/(FRACTION);
31     }
32     else{
33         return (x - (FRACTION/2))/(FRACTION);
34     }
35 }
36
37 int add_fp_fp(int x, int y){
38     // return fp + fp
39     return x + y;
40 }
41
42 int add_fp_int(int x, int n){
43     // return fp + int
44     return x + n * (FRACTION);
45 }
46
47 int sub_fp_fp(int x, int y){
48     // return fp - fp
49     return x - y;
50 }
51
52 int sub_fp_int(int x, int n){
53     // return fp - int
54     return x - n * (FRACTION);
55 }
56
57 int mul_fp_fp(int x, int y){
58     // return fp * fp
59     return ((int64_t)x) * y / (FRACTION);
60 }
61
62 int mul_fp_int(int x, int n){
63     // return fp * int
64     return x * n;
65 }
66
67 int div_fp_fp(int x, int y){
68     // return fp / fp
69     return ((int64_t)x) * (FRACTION) / y;
70 }
71
72 int div_fp_n(int x, int n){
73     // return fp / int
74     return x / n;
75 }

```

다음으로 load\_avg, recent\_cpu, priority의 계산을 위한 함수들을 구현했다. 아래 사진에 보이듯이 mlfqs\_calc\_priority(), mlfqs\_increment\_recent\_cpu(), mlfqs\_recalc\_per\_Sec(), mlfqs\_recalc\_all\_priority()이렇게 총 4개의 함수를 구현했다. 우선 design report에 작성했던 것처럼 모든 스레드의 priority는 매 4tick마다 recalculate 된다. 그리고 load\_avg와 모든 스레드의 recent\_cpu는 timer\_ticks()%TIMER\_FREQ==0, 즉 1초마다 업데이트 된다. timer\_interrupt()에서 이를 간편하게 실행시키기 위해서 두개의 함수 mlfqs\_recalc\_all\_priority()와 mlfqs\_recalc\_per\_Sec()로 분리해서 구현했다. mlfqs\_calc\_priority()에서는 모든 스레드의 priority를 재계산해주기 위해서 all\_list를 순회하면서 mlfqs\_calc\_priority()를 호출하도록 구현했다. mlfqs\_calc\_priority()는 thread\_set\_nice()에서 priority의 재계산을 할 때 사용하기 위해서 별도의 함수로 분리했다. 다음으로 매 timer interrupt마다 current thread의 recent\_cpu를 증가시켜 주어야 하기 때문에 mlfqs\_increment\_recent\_cpu()를 구현했다. 이 4개의 함수에는 "fixed\_point.h"에 구현된 함수들이 사용되었고, 계산 공식은 design report에도 적어 두었던 것처럼 APPENDIX에서 제공된 계산식을 따랐다.

```

220 void mlfqs_calc_priority (struct thread *t){
221     if (t == idle_thread) return;
222     t->priority = convert_to_int_round_nearest(sub_fp_int(sub_fp_fp(convert_to_fp(PRI_MAX),(div_fp_n(t->recent_cpu,4))), (t->nice*2)));
223     return;
224 }
225
226 void mlfqs_increment_recent_cpu (void){
227     if (thread_current() == idle_thread) return;
228     struct thread *t = thread_current ();
229     t->recent_cpu = add_fp_int(t->recent_cpu, 1);
230     return;
231 }
232
233 void mlfqs_recalc_per_Sec (void){
234     struct list_elem *e;
235     struct thread *t;
236     int load_avg_two;
237
238     // load_avg 계산
239     int n = list_size(&ready_list);
240     if (thread_current() != idle_thread) n++;
241     load_avg = add_fp_fp(mul_fp_fp(div_fp_fp(59<<14, 60<<14), load_avg), mul_fp_int(div_fp_fp(1<<14, 60<<14), n));
242     if (load_avg < 0) load_avg = 0;
243
244     // recent_cpu 계산
245     for (e = list_begin (&all_list); e != list_end (&all_list); e = list_next (e)){
246         t = list_entry(e, struct thread, allelem);
247         if (t == idle_thread) continue;
248         t->recent_cpu = add_fp_int(mul_fp_fp(div_fp_fp(mul_fp_int(load_avg, 2), add_fp_int(mul_fp_int(load_avg, 2), 1)), t->recent_cpu), t->nice);
249     }
250 }
251
252 void mlfqs_recalc_all_priority (void){
253     struct list_elem *e;
254     struct thread *t;
255     for (e = list_begin (&all_list); e != list_end (&all_list); e = list_next (e)){
256         t = list_entry(e, struct thread, allelem);
257         mlfqs_calc_priority(t);
258     }
259 }
260 /*구현한 부분 END*/

```

그리고 나서 timer\_interrupt()에서 mlfqs함수들을 이용해서 적절한 타이밍에 값들을 재계산하도록 구현했다.

```

169 /* Timer interrupt handler. */
170 static void
171 timer_interrupt (struct intr_frame *args UNUSED)
172 {
173     ticks++;
174     thread_tick ();
175     if (thread_mlfqs){
176         mlfqs_increment_recent_cpu();
177         if (timer_ticks() % 4 == 0){
178             mlfqs_recalc_all_priority();
179         }
180         if (timer_ticks() % TIMER_FREQ == 0){
181             mlfqs_recalc_per_Sec();
182         }
183     }
184     if(get_first_time_to_wakeup() <= ticks)
185         thread_wakeup(ticks);
186 }
187

```

다음으로는 구현되지 않았던 4개의 set/get 함수들을 구현했다. get함수들은 project에서 요구하는 형식에 맞추어서(load\_avg와 recent\_cpu의 경우에는 100을 곱해서 return) 구현했다. 그리고 thread\_set\_nice()의 경우에, APPENDIX에서 요구된 것처럼 nice값을 설정한 다음 바로 mlfqs\_calc\_priority()를 이용해서 priority를 재계산 해준다음 need\_yield()를 호출해서 yield해야하는지 검사하도록 구현했다.

```

527  /* Sets the current thread's nice value to NICE. */
528  void
529  thread_set_nice (int nice)
530  {
531      enum intr_level old_level = intr_disable();
532      struct thread *cur = thread_current ();
533      cur->nice = nice;
534      mlfqs_calc_priority(cur);
535      need_yield();
536      intr_set_level(old_level);
537  }
538
539  /* Returns the current thread's nice value. */
540  int
541  thread_get_nice (void)
542  {
543      enum intr_level old_level = intr_disable();
544      int value = thread_current ()->nice;
545      intr_set_level(old_level);
546      return value;
547  }
548
549  /* Returns 100 times the system load average. */
550  int
551  thread_get_load_avg (void)
552  {
553      enum intr_level old_level = intr_disable();
554      int value = convert_to_int_round_nearest(mul_fp_int(load_avg, 100));
555      intr_set_level(old_level);
556      return value;
557  }
558
559  /* Returns 100 times the current thread's recent_cpu value. */
560  int
561  thread_get_recent_cpu (void)
562  {
563      enum intr_level old_level = intr_disable();
564      int value = convert_to_int_round_nearest(mul_fp_int(thread_current()->recent_cpu, 100));
565      intr_set_level(old_level);
566      return value;
567  }

```

마지막으로 mlfqs가 실행되면 thread\_mlfqs값이 true로 설정된다. thread\_mlfqs를 이용해서 thread\_set\_priority(), lock\_acquire(), lock\_release() 세 함수에서 priority donation 관련 코드들이 동작하지 않도록 구현했다.

```

if (thread_mlfqs) return;

```



## 5. Result Verification

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhdtaa1
1
LLooaaddiinngg.....
Kernel command line: -q run alarm-single
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 392,806,400 loops/s.
Boot complete.
Executing 'alarm-single':
(alarm-single) begin
(alarm-single) Creating 5 threads to sleep 1 times each.
(alarm-single) Thread 0 sleeps 10 ticks each time,
(alarm-single) thread 1 sleeps 20 ticks each time, and so on.
(alarm-single) If successful, product of iteration count and
(alarm-single) sleep duration will appear in nondescending order.
(alarm-single) thread 0: duration=10, iteration=1, product=10
(alarm-single) thread 1: duration=20, iteration=1, product=20
(alarm-single) thread 2: duration=30, iteration=1, product=30
(alarm-single) thread 3: duration=40, iteration=1, product=40
(alarm-single) thread 4: duration=50, iteration=1, product=50
(alarm-single) end
Execution of 'alarm-single' complete.
Timer: 275 ticks
Thread: 250 idle ticks, 25 kernel ticks, 0 user ticks
Console: 987 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhdtaa1
1
LLooaaddiinngg.....
Kernel command line: -q run alarm-simultaneous
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 396,492,800 loops/s.
Boot complete.
Executing 'alarm-simultaneous':
(alarm-simultaneous) begin
(alarm-simultaneous) Creating 3 threads to sleep 5 times each.
(alarm-simultaneous) Each thread sleeps 10 ticks each time.
(alarm-simultaneous) Within an iteration, all threads should wake up on the same tick.
(alarm-simultaneous) iteration 0, thread 0: woke up after 10 ticks
(alarm-simultaneous) iteration 0, thread 1: woke up 0 ticks later
(alarm-simultaneous) iteration 0, thread 2: woke up 0 ticks later
(alarm-simultaneous) iteration 1, thread 0: woke up 10 ticks later
(alarm-simultaneous) iteration 1, thread 1: woke up 0 ticks later
(alarm-simultaneous) iteration 1, thread 2: woke up 0 ticks later
(alarm-simultaneous) iteration 2, thread 0: woke up 10 ticks later
(alarm-simultaneous) iteration 2, thread 1: woke up 0 ticks later
(alarm-simultaneous) iteration 2, thread 2: woke up 0 ticks later
(alarm-simultaneous) iteration 3, thread 0: woke up 10 ticks later
(alarm-simultaneous) iteration 3, thread 1: woke up 0 ticks later
(alarm-simultaneous) iteration 3, thread 2: woke up 0 ticks later
(alarm-simultaneous) iteration 4, thread 0: woke up 10 ticks later
(alarm-simultaneous) iteration 4, thread 1: woke up 0 ticks later
(alarm-simultaneous) iteration 4, thread 2: woke up 0 ticks later
(alarm-simultaneous) end
Execution of 'alarm-simultaneous' complete.
Timer: 281 ticks
Thread: 250 idle ticks, 31 kernel ticks, 0 user ticks
Console: 1615 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
Executing 'alarm-multiple':
(alarm-multiple) begin
(alarm-multiple) Creating 5 threads to sleep 7 times each.
(alarm-multiple) Thread 0 sleeps 10 ticks each time,
(alarm-multiple) thread 1 sleeps 20 ticks each time, and so on.
(alarm-multiple) If successful, product of iteration count and
(alarm-multiple) sleep duration will appear in nondescending order.
(alarm-multiple) thread 0: duration=10, iteration=1, product=10
(alarm-multiple) thread 1: duration=20, iteration=1, product=20
(alarm-multiple) thread 0: duration=10, iteration=2, product=20
(alarm-multiple) thread 2: duration=30, iteration=1, product=30
(alarm-multiple) thread 0: duration=10, iteration=3, product=30
(alarm-multiple) thread 3: duration=40, iteration=1, product=40
(alarm-multiple) thread 1: duration=20, iteration=2, product=40
(alarm-multiple) thread 0: duration=10, iteration=4, product=40
(alarm-multiple) thread 4: duration=50, iteration=1, product=50
(alarm-multiple) thread 0: duration=10, iteration=5, product=50
(alarm-multiple) thread 2: duration=30, iteration=2, product=60
(alarm-multiple) thread 1: duration=20, iteration=3, product=60
(alarm-multiple) thread 0: duration=10, iteration=6, product=60
(alarm-multiple) thread 0: duration=10, iteration=7, product=70
(alarm-multiple) thread 3: duration=40, iteration=2, product=80
(alarm-multiple) thread 1: duration=20, iteration=4, product=80
(alarm-multiple) thread 2: duration=30, iteration=3, product=90
(alarm-multiple) thread 4: duration=50, iteration=2, product=100
(alarm-multiple) thread 1: duration=20, iteration=5, product=100
(alarm-multiple) thread 3: duration=40, iteration=3, product=120
(alarm-multiple) thread 2: duration=30, iteration=4, product=120
(alarm-multiple) thread 1: duration=20, iteration=6, product=120
(alarm-multiple) thread 1: duration=20, iteration=7, product=140
(alarm-multiple) thread 4: duration=50, iteration=3, product=150
(alarm-multiple) thread 2: duration=30, iteration=5, product=150
(alarm-multiple) thread 3: duration=40, iteration=4, product=160
(alarm-multiple) thread 2: duration=30, iteration=6, product=180
(alarm-multiple) thread 4: duration=50, iteration=4, product=200
(alarm-multiple) thread 3: duration=40, iteration=5, product=200
(alarm-multiple) thread 2: duration=30, iteration=7, product=210
(alarm-multiple) thread 3: duration=40, iteration=6, product=240
(alarm-multiple) thread 4: duration=50, iteration=5, product=250
(alarm-multiple) thread 3: duration=40, iteration=7, product=280
(alarm-multiple) thread 4: duration=50, iteration=6, product=300
(alarm-multiple) thread 4: duration=50, iteration=7, product=350
(alarm-multiple) end
Execution of 'alarm-multiple' complete.
Timer: 584 ticks
Thread: 550 idle ticks, 34 kernel ticks, 0 user ticks
Console: 2955 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhdtaa1
1
LLooaaddiinngg.....
Kernel command line: -q run alarm-priority
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 396,083,200 loops/s.
Boot complete.
Executing 'alarm-priority':
(alarm-priority) begin
(alarm-priority) Thread priority 30 woke up.
(alarm-priority) Thread priority 29 woke up.
(alarm-priority) Thread priority 28 woke up.
(alarm-priority) Thread priority 27 woke up.
(alarm-priority) Thread priority 26 woke up.
(alarm-priority) Thread priority 25 woke up.
(alarm-priority) Thread priority 24 woke up.
(alarm-priority) Thread priority 23 woke up.
(alarm-priority) Thread priority 22 woke up.
(alarm-priority) Thread priority 21 woke up.
(alarm-priority) end
Execution of 'alarm-priority' complete.
Timer: 529 ticks
Thread: 490 idle ticks, 39 kernel ticks, 0 user ticks
Console: 840 characters output
Keyboard: 0 keys pressed
Powering off...
```



```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q run alarm-zero
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 398,950,400 loops/s.
Boot complete.
Executing 'alarm-zero':
(alarm-zero) begin
(alarm-zero) PASS
(alarm-zero) end
Execution of 'alarm-zero' complete.
Timer: 26 ticks
Thread: 1 idle ticks, 25 kernel ticks, 0 user ticks
Console: 385 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q run priority-change
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 389,120,000 loops/s.
Boot complete.
Executing 'priority-change':
(priority-change) begin
(priority-change) Creating a high-priority thread 2.
(priority-change) Thread 2 now lowering priority.
(priority-change) Thread 2 should have just lowered its priority.
(priority-change) Thread 2 exiting.
(priority-change) Thread 2 should have just exited.
(priority-change) end
Execution of 'priority-change' complete.
Timer: 26 ticks
Thread: 0 idle ticks, 26 kernel ticks, 0 user ticks
Console: 649 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q run priority-donate-multiple
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 398,950,400 loops/s.
Boot complete.
Executing 'priority-donate-multiple':
(priority-donate-multiple) begin
(priority-donate-multiple) Main thread should have priority 32. Actual priority: 32.
(priority-donate-multiple) Main thread should have priority 33. Actual priority: 33.
(priority-donate-multiple) Thread b acquired lock b.
(priority-donate-multiple) Thread b finished.
(priority-donate-multiple) Thread b should have just finished.
(priority-donate-multiple) Main thread should have priority 32. Actual priority: 32.
(priority-donate-multiple) Thread a acquired lock a.
(priority-donate-multiple) Thread a finished.
(priority-donate-multiple) Thread a should have just finished.
(priority-donate-multiple) Main thread should have priority 31. Actual priority: 31.
(priority-donate-multiple) end
Execution of 'priority-donate-multiple' complete.
Timer: 28 ticks
Thread: 0 idle ticks, 28 kernel ticks, 0 user ticks
Console: 1105 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q run alarm-negative
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 391,987,200 loops/s.
Boot complete.
Executing 'alarm-negative':
(alarm-negative) begin
(alarm-negative) PASS
(alarm-negative) end
Execution of 'alarm-negative' complete.
Timer: 26 ticks
Thread: 1 idle ticks, 25 kernel ticks, 0 user ticks
Console: 409 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q run priority-donate-one
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 392,806,400 loops/s.
Boot complete.
Executing 'priority-donate-one':
(priority-donate-one) begin
(priority-donate-one) This thread should have priority 32. Actual priority: 32.
(priority-donate-one) This thread should have priority 33. Actual priority: 33.
(priority-donate-one) acquire2: got the lock
(priority-donate-one) acquire2: done
(priority-donate-one) acquire1: got the lock
(priority-donate-one) acquire1: done
(priority-donate-one) acquire2, acquire1 must already have finished, in that order.
(priority-donate-one) This should be the last line before finishing this test.
(priority-donate-one) end
Execution of 'priority-donate-one' complete.
Timer: 26 ticks
Thread: 0 idle ticks, 26 kernel ticks, 0 user ticks
Console: 901 characters output
Keyboard: 0 keys pressed
Powering off...
```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q run priority-donate-multiple2
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 399,769,600 loops/s.
Boot complete.
Executing 'priority-donate-multiple2':
(priority-donate-multiple2) begin
(priority-donate-multiple2) Main thread should have priority 34. Actual priority: 34.
(priority-donate-multiple2) Main thread should have priority 36. Actual priority: 36.
(priority-donate-multiple2) Main thread should have priority 36. Actual priority: 36.
(priority-donate-multiple2) Thread b acquired lock b.
(priority-donate-multiple2) Thread b finished.
(priority-donate-multiple2) Thread a acquired lock a.
(priority-donate-multiple2) Thread a finished.
(priority-donate-multiple2) Thread c finished.
(priority-donate-multiple2) Threads b, a, c should have just finished, in that order.
(priority-donate-multiple2) Main thread should have priority 31. Actual priority: 31.
(priority-donate-multiple2) end
Execution of 'priority-donate-multiple2' complete.
Timer: 30 ticks
Thread: 0 idle ticks, 30 kernel ticks, 0 user ticks
Console: 1125 characters output
Keyboard: 0 keys pressed
Powering off...
```



```

SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PillLoo hddaa1
1
Llooaaddiinngg.....
Kernel command line: -q run priority-preempt
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 392,806,400 loops/s.
Boot complete.
Executing 'priority-preempt':
(priority-preempt) begin
(priority-preempt) Thread high-priority iteration 0
(priority-preempt) Thread high-priority iteration 1
(priority-preempt) Thread high-priority iteration 2
(priority-preempt) Thread high-priority iteration 3
(priority-preempt) Thread high-priority iteration 4
(priority-preempt) Thread high-priority done!
(priority-preempt) The high-priority thread should have already completed.
(priority-preempt) end
Execution of 'priority-preempt' complete.
Timer: 26 ticks
Thread: 0 idle ticks, 26 kernel ticks, 0 user ticks
Console: 778 characters output
Keyboard: 0 keys pressed
Powering off...

Booting from Hard Disk...
PillLoo hddaa1
1
Llooaaddiinngg.....
Kernel command line: -q run priority-condvar
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 394,444,800 loops/s.
Boot complete.
Executing 'priority-condvar':
(priority-condvar) begin
(priority-condvar) Thread priority 23 starting.
(priority-condvar) Thread priority 22 starting.
(priority-condvar) Thread priority 21 starting.
(priority-condvar) Thread priority 30 starting.
(priority-condvar) Thread priority 29 starting.
(priority-condvar) Thread priority 28 starting.
(priority-condvar) Thread priority 27 starting.
(priority-condvar) Thread priority 26 starting.
(priority-condvar) Thread priority 25 starting.
(priority-condvar) Thread priority 24 starting.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 30 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 29 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 28 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 27 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 26 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 25 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 24 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 23 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 22 woke up.
(priority-condvar) Signaling...
(priority-condvar) Thread priority 21 woke up.
(priority-condvar) end
Execution of 'priority-condvar' complete.
Timer: 33 ticks
Thread: 0 idle ticks, 33 kernel ticks, 0 user ticks
Console: 1667 characters output
Keyboard: 0 keys pressed
Powering off...

```

```

SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PillLoo hddaa1
1
Llooaaddiinngg.....
Kernel command line: -q run priority-sema
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 398,131,200 loops/s.
Boot complete.
Executing 'priority-sema':
(priority-sema) begin
(priority-sema) Thread priority 30 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 29 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 28 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 27 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 26 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 25 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 24 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 23 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 22 woke up.
(priority-sema) Back in main thread.
(priority-sema) Thread priority 21 woke up.
(priority-sema) Back in main thread.
(priority-sema) end
Execution of 'priority-sema' complete.
Timer: 29 ticks
Thread: 0 idle ticks, 29 kernel ticks, 0 user ticks
Console: 1192 characters output
Keyboard: 0 keys pressed
Powering off...

367 pages available in user pool.
Calibrating timer... 396,083,200 loops/s.
Boot complete.
Executing 'priority-donate-chain':
(priority-donate-chain) begin
(priority-donate-chain) main got lock.
(priority-donate-chain) main should have priority 3. Actual priority: 3.
(priority-donate-chain) main should have priority 6. Actual priority: 6.
(priority-donate-chain) main should have priority 9. Actual priority: 9.
(priority-donate-chain) main should have priority 12. Actual priority: 12.
(priority-donate-chain) main should have priority 15. Actual priority: 15.
(priority-donate-chain) main should have priority 18. Actual priority: 18.
(priority-donate-chain) main should have priority 21. Actual priority: 21.
(priority-donate-chain) thread 1 got lock
(priority-donate-chain) thread 1 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 2 got lock
(priority-donate-chain) thread 2 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 3 got lock
(priority-donate-chain) thread 3 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 4 got lock
(priority-donate-chain) thread 4 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 5 got lock
(priority-donate-chain) thread 5 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 6 got lock
(priority-donate-chain) thread 6 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 7 got lock
(priority-donate-chain) thread 7 should have priority 21. Actual priority: 21
(priority-donate-chain) thread 7 finishing with priority 21.
(priority-donate-chain) interloper 7 finished.
(priority-donate-chain) thread 6 finishing with priority 18.
(priority-donate-chain) interloper 6 finished.
(priority-donate-chain) thread 5 finishing with priority 15.
(priority-donate-chain) interloper 5 finished.
(priority-donate-chain) thread 4 finishing with priority 12.
(priority-donate-chain) interloper 4 finished.
(priority-donate-chain) thread 3 finishing with priority 9.
(priority-donate-chain) interloper 3 finished.
(priority-donate-chain) thread 2 finishing with priority 6.
(priority-donate-chain) interloper 2 finished.
(priority-donate-chain) thread 1 finishing with priority 3.
(priority-donate-chain) interloper 1 finished.
(priority-donate-chain) main finishing with priority 0.
(priority-donate-chain) end
Execution of 'priority-donate-chain' complete.
Timer: 35 ticks
Thread: 0 idle ticks, 35 kernel ticks, 0 user ticks
Console: 2636 characters output
Keyboard: 0 keys pressed
Powering off...

```



```

SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PiiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q -mlfqs run mlfqs-load-1
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 396,902,400 loops/s.
Boot complete.
Executing 'mlfqs-load-1':
(mlfqs-load-1) begin
(mlfqs-load-1) spinning for up to 45 seconds, please wait...
(mlfqs-load-1) load average rose to 0.5 after 42 seconds
(mlfqs-load-1) sleeping for another 10 seconds, please wait...
(mlfqs-load-1) load average fell back below 0.5 (to 0.43)
(mlfqs-load-1) PASS
(mlfqs-load-1) end
Execution of 'mlfqs-load-1' complete.
Timer: 5300 ticks
Thread: 1000 idle ticks, 4300 kernel ticks, 0 user ticks
Console: 650 characters output
Keyboard: 0 keys pressed
Powering off...

```

```

SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PiiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q -mlfqs run mlfqs-load-60
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 375,193,600 loops/s.
Boot complete.
Executing 'mlfqs-load-60':
(mlfqs-load-60) begin
(mlfqs-load-60) Starting 60 niced load threads...
(mlfqs-load-60) Starting threads took 0 seconds.
(mlfqs-load-60) After 0 seconds, load average=0.00.
(mlfqs-load-60) After 2 seconds, load average=1.98.
(mlfqs-load-60) After 4 seconds, load average=3.90.
(mlfqs-load-60) After 6 seconds, load average=5.75.
(mlfqs-load-60) After 8 seconds, load average=7.55.
(mlfqs-load-60) After 10 seconds, load average=9.28.
(mlfqs-load-60) After 12 seconds, load average=10.95.
(mlfqs-load-60) After 14 seconds, load average=12.57.
(mlfqs-load-60) After 16 seconds, load average=14.14.
(mlfqs-load-60) After 18 seconds, load average=15.65.
(mlfqs-load-60) After 20 seconds, load average=17.12.
(mlfqs-load-60) After 22 seconds, load average=18.53.
(mlfqs-load-60) After 24 seconds, load average=19.90.
(mlfqs-load-60) After 26 seconds, load average=21.22.
(mlfqs-load-60) After 28 seconds, load average=22.50.
(mlfqs-load-60) After 30 seconds, load average=23.74.
(mlfqs-load-60) After 32 seconds, load average=24.93.
(mlfqs-load-60) After 34 seconds, load average=26.09.
(mlfqs-load-60) After 36 seconds, load average=27.21.
(mlfqs-load-60) After 38 seconds, load average=28.29.
(mlfqs-load-60) After 40 seconds, load average=29.33.
(mlfqs-load-60) After 42 seconds, load average=30.34.
(mlfqs-load-60) After 44 seconds, load average=31.32.
(mlfqs-load-60) After 46 seconds, load average=32.26.
(mlfqs-load-60) After 48 seconds, load average=33.17.
(mlfqs-load-60) After 50 seconds, load average=34.06.
(mlfqs-load-60) After 52 seconds, load average=34.91.
(mlfqs-load-60) After 54 seconds, load average=35.73.
(mlfqs-load-60) After 56 seconds, load average=36.53.
(mlfqs-load-60) After 58 seconds, load average=37.30.
(mlfqs-load-60) After 60 seconds, load average=38.05.
(mlfqs-load-60) After 62 seconds, load average=36.79.
(mlfqs-load-60) After 64 seconds, load average=35.57.
(mlfqs-load-60) After 66 seconds, load average=34.39.
(mlfqs-load-60) After 68 seconds, load average=33.25.
(mlfqs-load-60) After 70 seconds, load average=32.14.

```

```

(mlfqs-load-60) After 96 seconds, load average=20.73.
(mlfqs-load-60) After 98 seconds, load average=20.04.
(mlfqs-load-60) After 100 seconds, load average=19.38.
(mlfqs-load-60) After 102 seconds, load average=18.74.
(mlfqs-load-60) After 104 seconds, load average=18.11.
(mlfqs-load-60) After 106 seconds, load average=17.51.
(mlfqs-load-60) After 108 seconds, load average=16.93.
(mlfqs-load-60) After 110 seconds, load average=16.37.
(mlfqs-load-60) After 112 seconds, load average=15.83.
(mlfqs-load-60) After 114 seconds, load average=15.30.
(mlfqs-load-60) After 116 seconds, load average=14.80.
(mlfqs-load-60) After 118 seconds, load average=14.30.
(mlfqs-load-60) After 120 seconds, load average=13.83.
(mlfqs-load-60) After 122 seconds, load average=13.37.
(mlfqs-load-60) After 124 seconds, load average=12.93.
(mlfqs-load-60) After 126 seconds, load average=12.50.
(mlfqs-load-60) After 128 seconds, load average=12.08.
(mlfqs-load-60) After 130 seconds, load average=11.68.
(mlfqs-load-60) After 132 seconds, load average=11.30.
(mlfqs-load-60) After 134 seconds, load average=10.92.
(mlfqs-load-60) After 136 seconds, load average=10.56.
(mlfqs-load-60) After 138 seconds, load average=10.21.
(mlfqs-load-60) After 140 seconds, load average=9.87.
(mlfqs-load-60) After 142 seconds, load average=9.54.
(mlfqs-load-60) After 144 seconds, load average=9.23.
(mlfqs-load-60) After 146 seconds, load average=8.92.
(mlfqs-load-60) After 148 seconds, load average=8.62.
(mlfqs-load-60) After 150 seconds, load average=8.34.
(mlfqs-load-60) After 152 seconds, load average=8.06.
(mlfqs-load-60) After 154 seconds, load average=7.79.
(mlfqs-load-60) After 156 seconds, load average=7.54.
(mlfqs-load-60) After 158 seconds, load average=7.29.
(mlfqs-load-60) After 160 seconds, load average=7.04.
(mlfqs-load-60) After 162 seconds, load average=6.81.
(mlfqs-load-60) After 164 seconds, load average=6.58.
(mlfqs-load-60) After 166 seconds, load average=6.37.
(mlfqs-load-60) After 168 seconds, load average=6.15.
(mlfqs-load-60) After 170 seconds, load average=5.95.
(mlfqs-load-60) After 172 seconds, load average=5.75.
(mlfqs-load-60) After 174 seconds, load average=5.56.
(mlfqs-load-60) After 176 seconds, load average=5.38.
(mlfqs-load-60) After 178 seconds, load average=5.20.
(mlfqs-load-60) end
Execution of 'mlfqs-load-60' complete.
Timer: 18827 ticks
Thread: 12799 idle ticks, 6028 kernel ticks, 0 user ticks
Console: 5366 characters output
Keyboard: 0 keys pressed
Powering off...

```

```
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPIiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q -mlfq run mlfqs-load-avg
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 398,540,800 loops/s.
Boot complete.
Executing 'mlfq-load-avg':
(mlfq-load-avg) begin
(mlfq-load-avg) Starting 60 load threads...
(mlfq-load-avg) Starting threads took 0 seconds.
(mlfq-load-avg) After 0 seconds, load average=0.00.
(mlfq-load-avg) After 2 seconds, load average=0.05.
(mlfq-load-avg) After 4 seconds, load average=0.16.
(mlfq-load-avg) After 6 seconds, load average=0.34.
(mlfq-load-avg) After 8 seconds, load average=0.58.
(mlfq-load-avg) After 10 seconds, load average=0.87.
(mlfq-load-avg) After 12 seconds, load average=1.22.
(mlfq-load-avg) After 14 seconds, load average=1.63.
(mlfq-load-avg) After 16 seconds, load average=2.09.
(mlfq-load-avg) After 18 seconds, load average=2.60.
(mlfq-load-avg) After 20 seconds, load average=3.15.
(mlfq-load-avg) After 22 seconds, load average=3.76.
(mlfq-load-avg) After 24 seconds, load average=4.41.
(mlfq-load-avg) After 26 seconds, load average=5.11.
(mlfq-load-avg) After 28 seconds, load average=5.85.
(mlfq-load-avg) After 30 seconds, load average=6.63.
(mlfq-load-avg) After 32 seconds, load average=7.45.
(mlfq-load-avg) After 34 seconds, load average=8.31.
(mlfq-load-avg) After 36 seconds, load average=9.21.
(mlfq-load-avg) After 38 seconds, load average=10.14.
(mlfq-load-avg) After 40 seconds, load average=11.11.
(mlfq-load-avg) After 42 seconds, load average=12.11.
(mlfq-load-avg) After 44 seconds, load average=13.15.
(mlfq-load-avg) After 46 seconds, load average=14.22.
(mlfq-load-avg) After 48 seconds, load average=15.32.
(mlfq-load-avg) After 50 seconds, load average=16.44.
(mlfq-load-avg) After 52 seconds, load average=17.60.
(mlfq-load-avg) After 54 seconds, load average=18.78.
(mlfq-load-avg) After 56 seconds, load average=20.00.
(mlfq-load-avg) After 58 seconds, load average=21.23.
(mlfq-load-avg) After 60 seconds, load average=22.49.
(mlfq-load-avg) After 62 seconds, load average=23.73.
(mlfq-load-avg) After 64 seconds, load average=24.86.
(mlfq-load-avg) After 66 seconds, load average=25.92.
(mlfq-load-avg) After 68 seconds, load average=26.86.
(mlfq-load-avg) After 70 seconds, load average=27.72.
(mlfq-load-avg) After 96 seconds, load average=30.72.
(mlfq-load-avg) After 98 seconds, load average=30.44.
(mlfq-load-avg) After 100 seconds, load average=30.11.
(mlfq-load-avg) After 102 seconds, load average=29.72.
(mlfq-load-avg) After 104 seconds, load average=29.31.
(mlfq-load-avg) After 106 seconds, load average=28.82.
(mlfq-load-avg) After 108 seconds, load average=28.29.
(mlfq-load-avg) After 110 seconds, load average=27.70.
(mlfq-load-avg) After 112 seconds, load average=27.06.
(mlfq-load-avg) After 114 seconds, load average=26.38.
(mlfq-load-avg) After 116 seconds, load average=25.65.
(mlfq-load-avg) After 118 seconds, load average=24.89.
(mlfq-load-avg) After 120 seconds, load average=24.08.
(mlfq-load-avg) After 122 seconds, load average=23.28.
(mlfq-load-avg) After 124 seconds, load average=22.51.
(mlfq-load-avg) After 126 seconds, load average=21.76.
(mlfq-load-avg) After 128 seconds, load average=21.04.
(mlfq-load-avg) After 130 seconds, load average=20.34.
(mlfq-load-avg) After 132 seconds, load average=19.67.
(mlfq-load-avg) After 134 seconds, load average=19.01.
(mlfq-load-avg) After 136 seconds, load average=18.38.
(mlfq-load-avg) After 138 seconds, load average=17.77.
(mlfq-load-avg) After 140 seconds, load average=17.18.
(mlfq-load-avg) After 142 seconds, load average=16.61.
(mlfq-load-avg) After 144 seconds, load average=16.06.
(mlfq-load-avg) After 146 seconds, load average=15.53.
(mlfq-load-avg) After 148 seconds, load average=15.01.
(mlfq-load-avg) After 150 seconds, load average=14.52.
(mlfq-load-avg) After 152 seconds, load average=14.03.
(mlfq-load-avg) After 154 seconds, load average=13.57.
(mlfq-load-avg) After 156 seconds, load average=13.12.
(mlfq-load-avg) After 158 seconds, load average=12.68.
(mlfq-load-avg) After 160 seconds, load average=12.26.
(mlfq-load-avg) After 162 seconds, load average=11.86.
(mlfq-load-avg) After 164 seconds, load average=11.46.
(mlfq-load-avg) After 166 seconds, load average=11.08.
(mlfq-load-avg) After 168 seconds, load average=10.71.
(mlfq-load-avg) After 170 seconds, load average=10.36.
(mlfq-load-avg) After 172 seconds, load average=10.02.
(mlfq-load-avg) After 174 seconds, load average=9.68.
(mlfq-load-avg) After 176 seconds, load average=9.36.
(mlfq-load-avg) After 178 seconds, load average=9.05.
(mlfq-load-avg) end
Execution of 'mlfq-load-avg' complete.
Timer: 18828 ticks
Thread: 6898 idle ticks, 11930 kernel ticks, 0 user ticks
Console: 5461 characters output
Keyboard: 0 keys pressed
Powering off...
```



SeaBIOS (version 1.13.0-1ubuntu1)

Booting from Hard Disk...

PPiLLoo hhddaa1

1

LLooaaddiinnngg.....

Kernel command line: -q -mlfqs run mlfqs-recent-1

Pintos booting with 3,968 kB RAM...

367 pages available in kernel pool.

367 pages available in user pool.

Calibrating timer... 397,312,000 loops/s.

Boot complete.

Executing 'mlfqs-recent-1':

(mlfqs-recent-1) begin

(mlfqs-recent-1) Sleeping 10 seconds to allow recent\_cpu to decay, please wait...

(mlfqs-recent-1) After 2 seconds, recent\_cpu is 7.39, load\_avg is 0.03.

(mlfqs-recent-1) After 4 seconds, recent\_cpu is 13.57, load\_avg is 0.06.

(mlfqs-recent-1) After 6 seconds, recent\_cpu is 19.57, load\_avg is 0.10.

(mlfqs-recent-1) After 8 seconds, recent\_cpu is 25.37, load\_avg is 0.13.

(mlfqs-recent-1) After 10 seconds, recent\_cpu is 30.99, load\_avg is 0.15.

(mlfqs-recent-1) After 12 seconds, recent\_cpu is 36.44, load\_avg is 0.18.

(mlfqs-recent-1) After 14 seconds, recent\_cpu is 41.72, load\_avg is 0.21.

(mlfqs-recent-1) After 16 seconds, recent\_cpu is 46.83, load\_avg is 0.24.

(mlfqs-recent-1) After 18 seconds, recent\_cpu is 51.79, load\_avg is 0.26.

(mlfqs-recent-1) After 20 seconds, recent\_cpu is 56.58, load\_avg is 0.28.

(mlfqs-recent-1) After 22 seconds, recent\_cpu is 61.22, load\_avg is 0.31.

(mlfqs-recent-1) After 24 seconds, recent\_cpu is 65.71, load\_avg is 0.33.

(mlfqs-recent-1) After 26 seconds, recent\_cpu is 70.07, load\_avg is 0.35.

(mlfqs-recent-1) After 28 seconds, recent\_cpu is 74.29, load\_avg is 0.37.

(mlfqs-recent-1) After 30 seconds, recent\_cpu is 78.37, load\_avg is 0.39.

(mlfqs-recent-1) After 32 seconds, recent\_cpu is 82.32, load\_avg is 0.41.

(mlfqs-recent-1) After 34 seconds, recent\_cpu is 86.15, load\_avg is 0.43.

(mlfqs-recent-1) After 36 seconds, recent\_cpu is 89.85, load\_avg is 0.45.

(mlfqs-recent-1) After 38 seconds, recent\_cpu is 93.43, load\_avg is 0.47.

(mlfqs-recent-1) After 40 seconds, recent\_cpu is 96.91, load\_avg is 0.49.

(mlfqs-recent-1) After 42 seconds, recent\_cpu is 100.26, load\_avg is 0.50.

(mlfqs-recent-1) After 44 seconds, recent\_cpu is 103.51, load\_avg is 0.52.

(mlfqs-recent-1) After 46 seconds, recent\_cpu is 106.66, load\_avg is 0.54.

(mlfqs-recent-1) After 48 seconds, recent\_cpu is 109.71, load\_avg is 0.55.

(mlfqs-recent-1) After 50 seconds, recent\_cpu is 112.65, load\_avg is 0.57.

(mlfqs-recent-1) After 52 seconds, recent\_cpu is 115.49, load\_avg is 0.58.

(mlfqs-recent-1) After 54 seconds, recent\_cpu is 118.24, load\_avg is 0.59.

(mlfqs-recent-1) After 56 seconds, recent\_cpu is 120.91, load\_avg is 0.61.

(mlfqs-recent-1) After 58 seconds, recent\_cpu is 123.50, load\_avg is 0.62.

(mlfqs-recent-1) After 60 seconds, recent\_cpu is 126.00, load\_avg is 0.63.

(mlfqs-recent-1) After 62 seconds, recent\_cpu is 128.42, load\_avg is 0.64.

(mlfqs-recent-1) After 64 seconds, recent\_cpu is 130.76, load\_avg is 0.66.

(mlfqs-recent-1) After 66 seconds, recent\_cpu is 133.03, load\_avg is 0.67.

(mlfqs-recent-1) After 68 seconds, recent\_cpu is 135.23, load\_avg is 0.68.

(mlfqs-recent-1) After 70 seconds, recent\_cpu is 137.35, load\_avg is 0.69.

(mlfqs-recent-1) After 72 seconds, recent\_cpu is 139.40, load\_avg is 0.70.

(mlfqs-recent-1) After 74 seconds, recent\_cpu is 141.39, load\_avg is 0.71.

(mlfqs-recent-1) After 98 seconds, recent\_cpu is 160.76, load\_avg is 0.80.  
(mlfqs-recent-1) After 100 seconds, recent\_cpu is 162.03, load\_avg is 0.81.  
(mlfqs-recent-1) After 102 seconds, recent\_cpu is 163.28, load\_avg is 0.82.  
(mlfqs-recent-1) After 104 seconds, recent\_cpu is 164.48, load\_avg is 0.82.  
(mlfqs-recent-1) After 106 seconds, recent\_cpu is 165.66, load\_avg is 0.83.  
(mlfqs-recent-1) After 108 seconds, recent\_cpu is 166.79, load\_avg is 0.83.  
(mlfqs-recent-1) After 110 seconds, recent\_cpu is 167.89, load\_avg is 0.84.  
(mlfqs-recent-1) After 112 seconds, recent\_cpu is 168.94, load\_avg is 0.84.  
(mlfqs-recent-1) After 114 seconds, recent\_cpu is 169.97, load\_avg is 0.85.  
(mlfqs-recent-1) After 116 seconds, recent\_cpu is 170.95, load\_avg is 0.85.  
(mlfqs-recent-1) After 118 seconds, recent\_cpu is 171.90, load\_avg is 0.86.  
(mlfqs-recent-1) After 120 seconds, recent\_cpu is 172.82, load\_avg is 0.86.  
(mlfqs-recent-1) After 122 seconds, recent\_cpu is 173.72, load\_avg is 0.87.  
(mlfqs-recent-1) After 124 seconds, recent\_cpu is 174.59, load\_avg is 0.87.  
(mlfqs-recent-1) After 126 seconds, recent\_cpu is 175.42, load\_avg is 0.88.  
(mlfqs-recent-1) After 128 seconds, recent\_cpu is 176.24, load\_avg is 0.88.  
(mlfqs-recent-1) After 130 seconds, recent\_cpu is 177.03, load\_avg is 0.88.  
(mlfqs-recent-1) After 132 seconds, recent\_cpu is 177.77, load\_avg is 0.89.  
(mlfqs-recent-1) After 134 seconds, recent\_cpu is 178.49, load\_avg is 0.89.  
(mlfqs-recent-1) After 136 seconds, recent\_cpu is 179.20, load\_avg is 0.89.  
(mlfqs-recent-1) After 138 seconds, recent\_cpu is 179.89, load\_avg is 0.90.  
(mlfqs-recent-1) After 140 seconds, recent\_cpu is 180.54, load\_avg is 0.90.  
(mlfqs-recent-1) After 142 seconds, recent\_cpu is 181.18, load\_avg is 0.90.  
(mlfqs-recent-1) After 144 seconds, recent\_cpu is 181.81, load\_avg is 0.91.  
(mlfqs-recent-1) After 146 seconds, recent\_cpu is 182.41, load\_avg is 0.91.  
(mlfqs-recent-1) After 148 seconds, recent\_cpu is 182.97, load\_avg is 0.91.  
(mlfqs-recent-1) After 150 seconds, recent\_cpu is 183.52, load\_avg is 0.92.  
(mlfqs-recent-1) After 152 seconds, recent\_cpu is 184.07, load\_avg is 0.92.  
(mlfqs-recent-1) After 154 seconds, recent\_cpu is 184.58, load\_avg is 0.92.  
(mlfqs-recent-1) After 156 seconds, recent\_cpu is 185.08, load\_avg is 0.92.  
(mlfqs-recent-1) After 158 seconds, recent\_cpu is 185.58, load\_avg is 0.93.  
(mlfqs-recent-1) After 160 seconds, recent\_cpu is 186.06, load\_avg is 0.93.  
(mlfqs-recent-1) After 162 seconds, recent\_cpu is 186.52, load\_avg is 0.93.  
(mlfqs-recent-1) After 164 seconds, recent\_cpu is 186.96, load\_avg is 0.93.  
(mlfqs-recent-1) After 166 seconds, recent\_cpu is 187.39, load\_avg is 0.93.  
(mlfqs-recent-1) After 168 seconds, recent\_cpu is 187.80, load\_avg is 0.94.  
(mlfqs-recent-1) After 170 seconds, recent\_cpu is 188.21, load\_avg is 0.94.  
(mlfqs-recent-1) After 172 seconds, recent\_cpu is 188.58, load\_avg is 0.94.  
(mlfqs-recent-1) After 174 seconds, recent\_cpu is 188.96, load\_avg is 0.94.  
(mlfqs-recent-1) After 176 seconds, recent\_cpu is 189.31, load\_avg is 0.94.  
(mlfqs-recent-1) After 178 seconds, recent\_cpu is 189.67, load\_avg is 0.95.  
(mlfqs-recent-1) After 180 seconds, recent\_cpu is 190.00, load\_avg is 0.95.  
(mlfqs-recent-1) end

Execution of 'mlfqs-recent-1' complete.

Timer: 19001 ticks

Thread: 975 idle ticks, 18026 kernel ticks, 0 user ticks

Console: 7250 characters output

Keyboard: 0 keys pressed

Powering off...

SeaBIOS (version 1.13.0-1ubuntu1)

Booting from Hard Disk...

PPiLLoo hhddaa1

1

LLooaaddiinnngg.....

Kernel command line: -q -mlfqs run mlfqs-fair-20

Pintos booting with 3,968 kB RAM...

367 pages available in kernel pool.

367 pages available in user pool.

Calibrating timer... 396,083,200 loops/s.

Boot complete.

Executing 'mlfqs-fair-20':

(mlfqs-fair-20) begin

(mlfqs-fair-20) Starting 20 threads...

(mlfqs-fair-20) Starting threads took 0 ticks.

(mlfqs-fair-20) Sleeping 40 seconds to let threads run, please wait...

(mlfqs-fair-20) Thread 0 received 153 ticks.

(mlfqs-fair-20) Thread 1 received 149 ticks.

(mlfqs-fair-20) Thread 2 received 149 ticks.

(mlfqs-fair-20) Thread 3 received 148 ticks.

(mlfqs-fair-20) Thread 4 received 148 ticks.

(mlfqs-fair-20) Thread 5 received 153 ticks.

(mlfqs-fair-20) Thread 6 received 152 ticks.

(mlfqs-fair-20) Thread 7 received 153 ticks.

(mlfqs-fair-20) Thread 8 received 152 ticks.

(mlfqs-fair-20) Thread 9 received 148 ticks.

(mlfqs-fair-20) Thread 10 received 149 ticks.

(mlfqs-fair-20) Thread 11 received 148 ticks.

(mlfqs-fair-20) Thread 12 received 149 ticks.

(mlfqs-fair-20) Thread 13 received 148 ticks.

(mlfqs-fair-20) Thread 14 received 152 ticks.

(mlfqs-fair-20) Thread 15 received 152 ticks.

(mlfqs-fair-20) Thread 16 received 152 ticks.

(mlfqs-fair-20) Thread 17 received 153 ticks.

(mlfqs-fair-20) Thread 18 received 152 ticks.

(mlfqs-fair-20) Thread 19 received 149 ticks.

(mlfqs-fair-20) end

Execution of 'mlfqs-fair-20' complete.

Timer: 4033 ticks

Thread: 1000 idle ticks, 3033 kernel ticks, 0 user ticks

Console: 1463 characters output

Keyboard: 0 keys pressed

Powering off...

SeaBIOS (version 1.13.0-1ubuntu1)

Booting from Hard Disk...

PPiLLoo hhddaa1

1

LLooaaddiinnngg.....

Kernel command line: -q -mlfqs run mlfqs-fair-2

Pintos booting with 3,968 kB RAM...

367 pages available in kernel pool.

367 pages available in user pool.

Calibrating timer... 399,769,600 loops/s.

Boot complete.

Executing 'mlfqs-fair-2':

(mlfqs-fair-2) begin

(mlfqs-fair-2) Starting 2 threads...

(mlfqs-fair-2) Starting threads took 0 ticks.

(mlfqs-fair-2) Sleeping 40 seconds to let threads run, please wait...

(mlfqs-fair-2) Thread 0 received 1500 ticks.

(mlfqs-fair-2) Thread 1 received 1500 ticks.

(mlfqs-fair-2) end

Execution of 'mlfqs-fair-2' complete.

Timer: 4028 ticks

Thread: 1000 idle ticks, 3028 kernel ticks, 0 user ticks

Console: 634 characters output

Keyboard: 0 keys pressed

Powering off...

```

SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q -mlfqs run mlfqs-nice-2
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 397,721,600 loops/s.
Boot complete.
Executing 'mlfqs-nice-2':
(mlfqs-nice-2) begin
(mlfqs-nice-2) Starting 2 threads...
(mlfqs-nice-2) Starting threads took 0 ticks.
(mlfqs-nice-2) Sleeping 40 seconds to let threads run, please wait...
(mlfqs-nice-2) Thread 0 received 1936 ticks.
(mlfqs-nice-2) Thread 1 received 1064 ticks.
(mlfqs-nice-2) end
Execution of 'mlfqs-nice-2' complete.
Timer: 4026 ticks
Thread: 1000 idle ticks, 3026 kernel ticks, 0 user ticks
Console: 634 characters output
Keyboard: 0 keys pressed
Powering off...
SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q -mlfqs run mlfqs-block
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 394,444,800 loops/s.
Boot complete.
Executing 'mlfqs-block':
(mlfqs-block) begin
(mlfqs-block) Main thread acquiring lock.
(mlfqs-block) Main thread creating block thread, sleeping 25 seconds...
(mlfqs-block) Block thread spinning for 20 seconds...
(mlfqs-block) Block thread acquiring lock...
(mlfqs-block) Main thread spinning for 5 seconds...
(mlfqs-block) Main thread releasing lock.
(mlfqs-block) ...got it.
(mlfqs-block) Block thread should have already acquired lock.
(mlfqs-block) end
Execution of 'mlfqs-block' complete.
Timer: 3027 ticks
Thread: 500 idle ticks, 2527 kernel ticks, 0 user ticks
Console: 779 characters output
Keyboard: 0 keys pressed
Powering off...

```

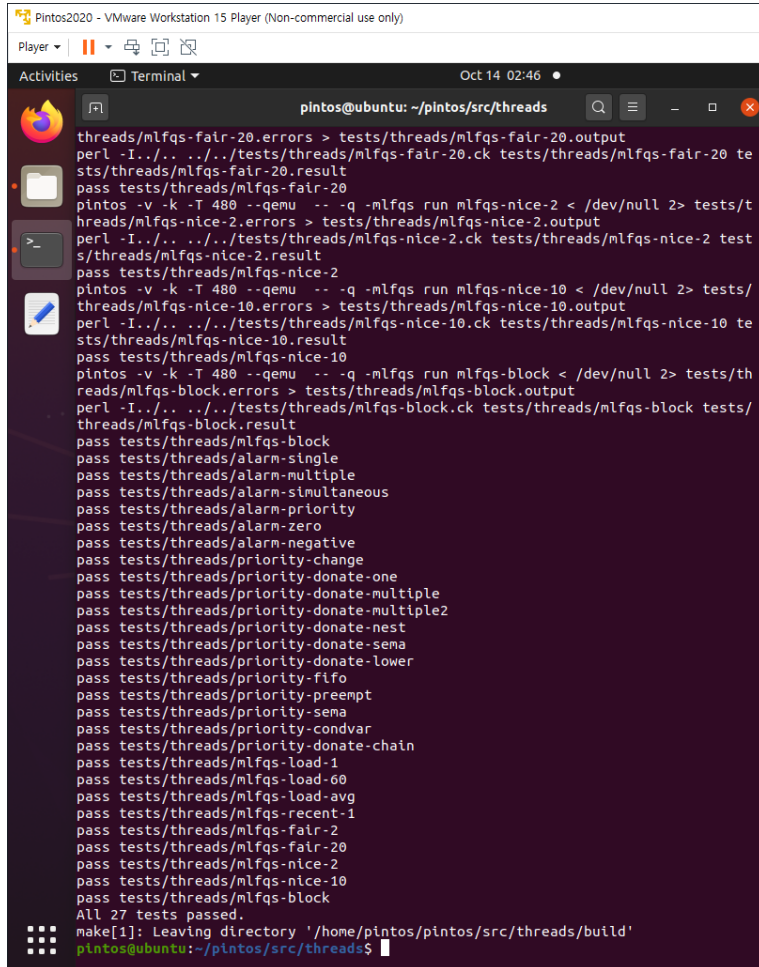
```

SeaBIOS (version 1.13.0-1ubuntu1)
Booting from Hard Disk...
PPiLLoo hhddaa1
1
LLOoaaddiinngg.....
Kernel command line: -q -mlfqs run mlfqs-nice-10
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 398,131,200 loops/s.
Boot complete.
Executing 'mlfqs-nice-10':
(mlfqs-nice-10) begin
(mlfqs-nice-10) Starting 10 threads...
(mlfqs-nice-10) Starting threads took 0 ticks.
(mlfqs-nice-10) Sleeping 40 seconds to let threads run, please wait...
(mlfqs-nice-10) Thread 0 received 688 ticks.
(mlfqs-nice-10) Thread 1 received 588 ticks.
(mlfqs-nice-10) Thread 2 received 505 ticks.
(mlfqs-nice-10) Thread 3 received 397 ticks.
(mlfqs-nice-10) Thread 4 received 317 ticks.
(mlfqs-nice-10) Thread 5 received 224 ticks.
(mlfqs-nice-10) Thread 6 received 149 ticks.
(mlfqs-nice-10) Thread 7 received 89 ticks.
(mlfqs-nice-10) Thread 8 received 41 ticks.
(mlfqs-nice-10) Thread 9 received 8 ticks.
(mlfqs-nice-10) end
Execution of 'mlfqs-nice-10' complete.
Timer: 4030 ticks
Thread: 1000 idle ticks, 3030 kernel ticks, 0 user ticks
Console: 999 characters output
Keyboard: 0 keys pressed
Powering off...

```

## 6. Conclusion

이번 과제의 목표인 Alarm Clock, Priority Scheduling, Advanced Scheduler를 모두 성공적으로 구현하였다.



```
Pintos2020 - VMware Workstation 15 Player (Non-commercial use only)
Player
Activities Terminal
pintos@ubuntu: ~/pintos/src/threads
threads/mlfqs-fair-20.errors > tests/threads/mlfqs-fair-20.output
perl -I../.. ../tests/threads/mlfqs-fair-20.ck tests/threads/mlfqs-fair-20 te
sts/threads/mlfqs-fair-20.result
pass tests/threads/mlfqs-fair-20
pintos -v -k -T 480 --qemu -- -q -mlfqs run mlfqs-nice-2 < /dev/null 2> tests/t
hreads/mlfqs-nice-2.errors > tests/threads/mlfqs-nice-2.output
perl -I../.. ../tests/threads/mlfqs-nice-2.ck tests/threads/mlfqs-nice-2 test
s/threads/mlfqs-nice-2.result
pass tests/threads/mlfqs-nice-2
pintos -v -k -T 480 --qemu -- -q -mlfqs run mlfqs-nice-10 < /dev/null 2> tests/t
hreads/mlfqs-nice-10.errors > tests/threads/mlfqs-nice-10.output
perl -I../.. ../tests/threads/mlfqs-nice-10.ck tests/threads/mlfqs-nice-10 te
sts/threads/mlfqs-nice-10.result
pass tests/threads/mlfqs-nice-10
pintos -v -k -T 480 --qemu -- -q -mlfqs run mlfqs-block < /dev/null 2> tests/th
reads/mlfqs-block.errors > tests/threads/mlfqs-block.output
perl -I../.. ../tests/threads/mlfqs-block.ck tests/threads/mlfqs-block tests/
threads/mlfqs-block.result
pass tests/threads/mlfqs-block
pass tests/threads/alarm-single
pass tests/threads/alarm-multiple
pass tests/threads/alarm-simultaneous
pass tests/threads/alarm-priority
pass tests/threads/alarm-zero
pass tests/threads/alarm-negative
pass tests/threads/priority-change
pass tests/threads/priority-donate-one
pass tests/threads/priority-donate-multiple
pass tests/threads/priority-donate-multiple2
pass tests/threads/priority-donate-nest
pass tests/threads/priority-donate-sema
pass tests/threads/priority-donate-lower
pass tests/threads/priority-fifo
pass tests/threads/priority-preempt
pass tests/threads/priority-sema
pass tests/threads/priority-condvar
pass tests/threads/priority-donate-chain
pass tests/threads/mlfqs-load-1
pass tests/threads/mlfqs-load-60
pass tests/threads/mlfqs-load-avg
pass tests/threads/mlfqs-recent-1
pass tests/threads/mlfqs-fair-2
pass tests/threads/mlfqs-fair-20
pass tests/threads/mlfqs-nice-2
pass tests/threads/mlfqs-nice-10
pass tests/threads/mlfqs-block
All 27 tests passed.
make[1]: Leaving directory '/home/pintos/pintos/src/threads/build'
pintos@ubuntu:~/pintos/src/threads$
```

이번 과제를 진행하면서 처음부터 제공된 함수들 간의 상관관계를 모두 파악해야 했고, 복잡한 함수의 흐름 속에서 새로운 개념(priority donation 과 같은)을 도입하여 구현해야해서 어려웠다. 하지만, 함수의 흐름을 그려보고 구현해야 할 것이 함수들 가운데서 어떻게 동작해야 하는지 논리적인 흐름을 따라 구현하니 마침내 구현에 성공하였다. 이번 과제를 통해 OS가 어떻게 thread들의 실행을 관리하는지, 여러 복잡한 관계의 thread들 가운데 priority donation을 사용하여 어떻게 이를 해결할 수 있는지, mlfqs scheduler를 사용하여 어떻게 thread들을 관리하는지 자세히 알 수 있는 시간이었다. pintoS는 실제로 우리가 사용하는 OS보다 간단한 OS인데도 직접 구현하려고 하니 이렇게 복잡한 것을 보니 실제 우리가 사용하고 있는 OS는 정확히 어떤 구조로 이루어져 있고 pintoS에 비해 추가적으로 어떤 것들이 더 필요한지 궁금해졌다.